



Available online at www.sciencedirect.com



INTEGRATION, the VLSI journal 38 (2004) 309–339

INTEGRATION
the VLSI journal

www.elsevier.com/locate/vlsi

Application of the object-oriented principles for hardware and embedded system design

R. Damaševičius*, V. Štuikys

Software Engineering Department, Kaunas University of Technology, Studentų 50, 51368-Kaunas, Lithuania

Received 12 February 2004; received in revised form 21 August 2004; accepted 24 August 2004

Abstract

As the complexity of hardware (HW) and embedded system design is constantly increasing, the researchers are seeking to develop new more abstract and productive design methods or adapt the existing ones from other domains such as software design. This paper addresses the problem of using the object-oriented (OO) design techniques in HW domain. The main OO design techniques are as follows: abstraction, separation of concerns, composition and generalization. The application of the OO design paradigm has many aspects: high-level specification of HW models using OO formal notations such as Petri Nets and UML diagrams, HW description using OO HW description languages such as VHDL extensions and SystemC, HW design using OO HW architectures, platforms and design patterns. In this paper, we present a comprehensive overview of the application of the OO design paradigm in HW and embedded system design domains and formulate its main principles, discuss the current achievements in the area, and outline the future trends.

© 2004 Elsevier B.V. All rights reserved.

Keywords: Object-oriented hardware design; Embedded system design; Modeling; UML; Class diagram.

1. Introduction

The usage of high-level models for specifying complex domain systems is a long-standing engineering tradition. Modeling helps to focus on the most important aspects of the design

*Corresponding author. Tel.: +370-37-300-399; fax: +370-37-300-352.

E-mail address: damarobe@soften.ktu.lt (R. Damaševičius).

problem, to communicate design ideas across the design group, as well as to validate the designed system and evaluate its characteristics before it is actually implemented. Thus, the development of high-level models based on a well-proven design methodology rather than design of the specific components (systems) for specific applications is becoming the primary concern of the designers.

Hardware (HW) domain is a particularly complex one. Modern chips can contain up to 25 M gates and may require 6–8 months to design [1]. The entire *Systems-on-Chip* (SoC) with multiple microprocessors, memory with embedded software (SW) and application-specific circuitry can be implemented on a single chip now. In fact, semiconductor technology currently allows much more complex systems than HW designers can actually design. The complexity of SoC in terms of logic transistors that can be integrated on a chip is increasing at the rate of 58% per year (*Moore's law*). However, the design productivity is increasing at the rate of 21% per year only. This fact is known in the EDA community as *design productivity gap* [2].

Most of current research efforts in the domain are directed at bridging this gap as well as raising the level of abstraction in design and unifying the HW and embedded SW design methodologies [3]. The researchers have to reconsider the existing HW modeling and design methodologies as well as to develop or adopt the new ones that can provide higher productivity and shorten time-to-market.

This paper analyzes the application of the object-oriented (OO) design (OOD) paradigm for HW and embedded system design. Note that we use the term 'design' as the HW community understands it, i.e., the development of HW models, components and systems. Though OOD of HW is not a new issue, and was discussed before [4–13], recently it has received an increased attention from researchers and designers. This interest is motivated by the fact that the introduction of the OOD methodology allowed to raise the level of abstraction and brought major productivity increase in SW design [14], and it is broadly expected to do the same in HW and embedded system design domains.

HW designers can increase flexibility and reusability of the modeled systems or their components by using the OOD techniques. The OOD techniques tend to increase the reuse of components, which results in higher-quality programs and faster system development. Since nowadays HW design is becoming more and more a programming activity similar to SW development, the application of the OOD techniques can be considered as a way to cope with these challenges.

The OOD paradigm assumes that systems are modeled using a set of classes and relationships between them. A class encapsulates the data and the operations applicable to the data. An instance of a class is an object. The system consists of the communicating objects.

An HW component can be perceived as an individual class with characteristics and operations, and decomposition of a system into objects is better defined and more explicit in HW than in SW. However, the application of the OOD methods in HW design domain so far had only a mixed success. Some concepts such as system decomposition into modules and information encapsulation and hiding were successfully adopted, whereas others (e.g., inheritance) were not so popular. The primary reason for this was the lack of the adequate abstractions (languages, metamodels) for expressing the OOD concepts in HW domain. Only recently, there has been a renewed interest in the application of OOD methods for HW design fueled by the arrival of SystemC [15] and adoption of UML [16] for embedded system design [17–22].

The aim of this paper is as follows: (1) To present an extensive survey of the application of the OOD methods and techniques in HW and embedded system design, including the OO HW architectures, high-level specification methods, and OO HW description languages (HDLs). (2) To discuss the application of design patterns for HW design. (3) To formulate the main principles of the OOD for HW and embedded system design domain. (4) To summarize the current achievements in the area and outline future trends.

Our recent contributions to the area of the OO HW design are also presented, including the high-level specification of HW models using UML class diagrams and the UML–VHDL metamodel for the automatic generation of VHDL components from the UML class diagrams.

The structure of the paper is as follows. Section 2 analyzes the application of the OOD techniques in HW design. Section 3 considers the application of design patterns for HW and embedded system design. Section 4 presents the main principles of the OOD paradigm adapted for HW and embedded system design domains. Section 5 illustrates the analyzed concepts in a case study. Section 6 evaluates and discusses current achievements in the domain. Finally, Section 7 presents the conclusions and outlines future trends.

2. Analysis of Approaches for OO HW Design

2.1. OO HW description languages

2.1.1. Extensions of VHDL

Here we consider the approaches that extend an existing HW description language (HDL) for OO HW design, only. A vast majority of these approaches uses VHDL as a base language. The reasons why VHDL is so often selected by the researchers to include support for the OOD concepts are as follows: (1) VHDL is an industry-wide standard for HW design [23]. (2) The language already has an adequate basis mechanism to separate the component interface from its implementation using entity and architecture constructs. (3) There are some similarities between VHDL structural descriptions and objects [24].

The problem is what abstractions should be used to extend VHDL in order to support the principles of OOD, and how to solve the related problems of modeling and synthesis. Basically, there are two approaches for introducing the OOD concepts into the domain by extending the component abstraction of VHDL or the type system as follows: the entity-based and type-based ones.

The *entity-based* approaches consider a VHDL entity as an abstraction of an HW object. The *entity* is extended with an additional interface for methods, or a new class construct based on the entity is introduced. The entity declaration is treated as a class definition, which specifies the interface of the HW object. The *architecture* is treated as an implementation of the interface, expressed in terms of *concurrent statements*, *processes*, and *component instances*. A new design entity can be defined by inheriting the interface and implementation from a parent entity. The new interface can be augmented with additional generic constants and ports, and the new architecture can be extended with additional declarations, processes, and component instances.

Thus, the entity-based approaches allow structural inheritance and can be used for writing an initial OO specification of HW and reusing parts of old specifications. The OO specification can be

simulated at a high level of abstraction, but its refinement (synthesis) to the lower levels of abstraction is complicated. The examples of the application of the entity-based approach are OO-VHDL [25], VHDL_OBJ [26], VHDL++ [27], and LaMI proposal [28].

The *type-based* approaches provide object-orientation based on an extension of the language-type system by a class type. Basically, such a class type is an abstract data type. It encapsulates its structure and models its behavior using procedures. Whereas the entity-based approaches focus on the hierarchical decomposition based on the functional units, the type-based approaches use the data-oriented decomposition. Various proposals differ with respect to their class concepts, the way classes can be instantiated, and the details of the inheritance mechanism and polymorphism.

Compared to the entity-based approach, the type-based approach is closer to the HW domain, because it allows a more simple translation into the HW-oriented data representations. Thus, the synthesis is simpler. However, HW components cannot be described using the types alone and should include functional units, too. The examples of the application of the type-based approach are Objective VHDL [29], SUAVE [30]. A more extensive survey of the languages that extend VHDL with the OO features can also be found in [31–33].

2.1.2. *SystemVerilog*

SystemVerilog [34] is a rich set of extensions to the IEEE 1364 Verilog-2001 HDL. One of these extensions adds the OO abstractions (*classes*) that are similar to C++. A class can contain data declarations (*properties*), tasks and functions for operating on the data (*methods*). The properties and methods together define the contents and capabilities of an *object*. Classes can have inheritance and public or private protection, as in C++. Objects can be dynamically created, deleted and assigned values. SystemVerilog classes are dynamic by nature and can be used for verification routines, highly abstract system-level modeling and test-bench modeling. Because of the dynamic nature of classes, they are not synthesizable.

These OO extensions address two major issues of the HDL-based design. First, the object-orientation allows modeling very large designs with concise, accurate and intuitive code. Second, high-level OO test programs allow to efficiently and effectively verify large designs. SystemVerilog achieves the improved specification of a design by allowing the related domain functionality to be described as a single object. It allows a single designer to be responsible for the specification and development of an object, rather than spreading that design knowledge diffusely throughout a system, which can aggravate design reuse and maintenance.

2.2. *Adaptation of an existing OO language to HW design*

Another trend in adopting the OOD concepts to HW design domain is based on using the existing OO programming language adopted from the SW design domain. Though approaches based on other OO programming languages such as Java exist (e.g., JavaSynth [35]), the main trend is focused on the adaptation of C++ (currently a de facto standard SW development language) for HW modeling and design.

Originally intended for SW design, C++ has no support required for accurately dealing with HW models. The problem is how to adapt C++ for HW design and to solve the synthesizability issue. C++ provides only two levels of abstraction: algorithmic and object, whereas an HW specification has four levels: algorithmic, modular, cycle-accurate, and register-transfer level

(RTL) [36]. The algorithmic level specifies only the behavior of the system, with no specific implementation detail. At the modular level, the system is partitioned into components that communicate through a clearly (although not necessarily completely) specified protocol. The cycle accuracy level introduces the notion of a clock and a time at which events occur. RTL specifies the implementation of the events without relying on a particular implementation technology.

Each level of abstraction requires the corresponding language support. Given that C++ is used to describe SW algorithms, the algorithmic level requires almost no extensions (except of a precisely sized HW-oriented data type). The object level of C++ can be used to describe the modular level of HW. The other levels of abstraction are not present in C++, and thus require introduction when designed for HW. According to [37], there are two approaches for building this type of HW support:

- (1) The *syntax extension* approach is based on adding keywords to the C/C++ language, in order to support HW description at a high level as a basis for synthesis (e.g., HardwareC [38], Matisse [39], SpecC [40]), and requires the development of separate compilers, simulators, and synthesis tools to manipulate with new syntax.
- (2) The *class libraries* approach is based on using class libraries that model the HW parts of the embedded system. This approach allows reusing the existing language development framework. As such, this is the most popular approach in the EDA community, now. Below, we analyze two HDLs based on the extension of C++ with class libraries for modeling HW: SystemC and SystemC-Plus.

SystemC [15,41,42] is a C++ class library and a methodology that allows to model embedded systems, including both HW and SW. Despite the fact that SystemC is build on top of C++, currently it does not encourage the usage of the OOD techniques for HW synthesis [43,44].

The important features of SystemC are briefly discussed below. A synthesizable HW model described in SystemC looks very similar to an equivalent VHDL model with only few syntactic differences. The basic block of a SystemC program is a *module*. A module is similar to the entity in VHDL. It is an abstract representation of a functional unit, without specifying any implementation details. Each module has a set of *ports* through which it interacts with the outside world. The ports can be input, output, or input/output ones. Individual modules communicate with one another through *signals* that connect the ports of the modules. The code that implements the algorithm of a module is encapsulated in one or more *processes*. All processes are executed concurrently, therefore a SystemC module class has in fact only one method.

Fig. 1 presents two simple SystemC modules and demonstrates the application of encapsulation, separation of data and methods, and inheritance. Fig. 1(a) shows the And gate, and (b) shows a derived class *Half_adder* that inherits the ports *X1*, *X2*, *Y* and the method *process1()* from the *And_gate*.

Unfortunately, the OOD paradigm cannot really be applied for HW design using SystemC [45]. Only encapsulation is fully supported by the SystemC module. Commonality can be expressed using abstract classes such as interfaces, and variation can be implemented using multiple inheritance. Inheritance of other SystemC modules is not supported for synthesis as well as polymorphism. Basically, SystemC partially supports the OO HW modeling, but not the OO HW design, because the commercially available synthesis tools for SystemC currently cannot handle

<pre> SC_MODULE (And_gate) { public: sc_in<sc_bit> X1, X2; sc_out<sc_bit> Y; void process_and() { Y = X1 && X2; // and } SC_CTOR(And_gate) { SC_METHOD(process_and); sensitive << X1 << X2; } };</pre> <p style="text-align: center;">(a)</p>	<pre> class Half_adder: public And_gate { public: sc_out<sc_bit> Z; SC_HAS_PROCESS(And_gate); void process_xor() { Y = X1 ^ X2; // xor } Half_adder(const sc_module_name& name): And_gate(name) { SC_METHOD(process_xor); sensitive << X1 << X2; } };</pre> <p style="text-align: center;">(b)</p>
---	--

Fig. 1. SystemC description of And gate (a) and Half adder (b).

the OO specifications. Thus, the designer has to manually rewrite the source code of external objects and translate certain high-level constructs like threads and FIFOs. However, there are efforts to extend SystemC as well, in order to provide more support for synthesizability of the OO specifications, such as SystemC-Plus [46].

SystemC-Plus is the OOD methodology based on SystemC that allows the usage of objects for HW modeling and synthesis. SystemC libraries are augmented with an additional class library and coding styles giving the ability for the designer to use the OO techniques (such as polymorphism) for HW modeling without losing the possibility of synthesizing these models. The extensions to common SystemC were necessary to avoid the usage of the non-synthesizable language constructs and to reflect some features of the HW model in high-level simulation.

SystemC-Plus supports the following OOD concepts: data encapsulation, inheritance, method-based communication, polymorphism, and classes. SystemC-Plus supports the synthesis of the following OO principles that are present in C++: class member functions, constructors, objects and arrays as data members, and inheritance from class templates, multiple inheritance and virtual base classes. Polymorphic objects and dynamic dispatching of virtual functions can be synthesized, too.

2.3. OO HW architectures

The implementation of the OO HW model is another subject of research. Every high-level HW specification eventually must be mapped into low-level (RTL) implementation (gates and wires). The problem is what HW architecture should be selected for an application, how it must be specified and implemented. Basically, there can be three approaches for implementing HW architectures in terms of OOD:

- (1) The *behavioral* model understands an object as an encapsulated data packet transported via wires (see Fig. 2) [35]. The objects are not treated as components, but rather as connections between components, which are themselves represented by the methods. A method corresponds to an HW component that processes the data packets and returns the result.



Fig. 2. Behavioral model of an object.

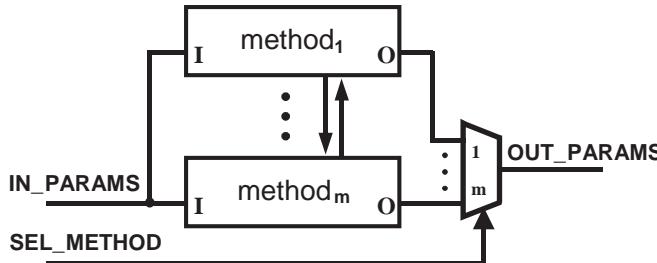


Fig. 3. Structural model of an object.

The object's state (Q, Q^+) is treated equally with other method's parameters that are supplied as inputs (I) and outputs (O) to the HW component. A return value of the method becomes an output of the component. This model supports encapsulation of the data and its separation from functionality.

- (2) The *structural* model represents *physical components* such as VHDL entities as objects (see Fig. 3) [24]. The object is an HW component that receives information about a method to be executed via wires. The objects communicate with one another via method calls, thus the methods are interpreted as ports/connections. Inheritance is supported by deriving a child entity from a parent entity with the addition of new methods. Encapsulation is supported by an additional interface to allow for the invocation of methods. Polymorphism is supported by dynamically activating different design entities during run-time.

Here, SEL_METHOD is a method identifier, IN_PARAMS are the inputs to the object's methods, and OUT_PARAMS are the outputs.

- (3) The *FSM-based* model represents each object as an FSM (see Fig. 4) [33]. The target architecture has a memory element for storing a state, state transition and output logic, and a feedback of the next state into the state memory. A method corresponds to an HW component that implements the functionality described by the methods. The object's state and the method's parameters are both inputs and outputs of the HW component since they may be read and modified. Some bits in the object storage can be used to show its class membership (i.e., type). Adding new methods to the entity and extending the object's state support inheritance. Polymorphism is achieved using a *switch-case* construct (implemented as a *mux*) to test the object's type at run-time and call the appropriate method. Message exchange between objects is also supported via a network of channels.

Here, SEL_METHOD is a method identifier, IN_PARAMS are the inputs to the object's methods, OUT_PARAMS are the outputs, DONE signals the end of a method's

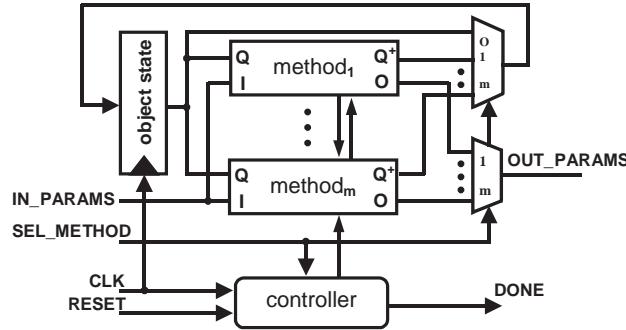


Fig. 4. FSM-based model of an object (according to [33]).

execution, a memory element stores the object's state, CLK is a clock signal, and RESET is a reset signal.

We summarize these three interpretations of objects in HW architectures as follows. The behavioral model is the simplest one: it allows encapsulation only. The structural model allows stateless static objects with inheritance. The FSM-based model allows static objects with support for inheritance, message passing and polymorphism. Dynamic objects cannot be implemented in HW directly, however, they can be emulated when HW objects are combined with SW objects in a shared memory. This approach is used in HW platforms considered in the next section.

2.4. OO HW Platforms

The OOD principles can also be applied for *platform-based design* [47] of embedded systems. In this case, platforms are modeled using a set of classes and relationships between them. We analyze the OO-ASIP and Enodia platforms below.

2.4.1. OO-ASIP platform

In the OO-ASIP [48] platform (see Fig. 5), the invocation of object's methods is used as an instruction to be executed on a processor specifying the object and additional arguments as instruction operands. This view results in identifying *public* methods of the class library as the instruction set of a corresponding processor. An OO-ASIP is a processor with its instruction set defined by selected methods of an “HW class library”, and with support for dynamic dispatch of virtual methods to both HW and SW. An embedded application is modeled using the “HW class library” and compiled to the instructions of the corresponding OO-ASIP. “HW class library” can incrementally evolve to model new applications, and thus ensures reuse in future applications. Such a processor is able to execute any program comprised of objects of that class hierarchy, i.e. all applications can be modeled with that HW class library.

The system designed on the OO-ASIP platform works as follows. The Method Invocation Unit (MIU) reads method invocation commands from the instruction memory and designates a method (*opcode*) and the called object (*operand(s)*). Based on the class membership of the called object, the MIU invokes the corresponding Functional Unit (FU), which implements the HW method, or dispatches a method call to the SW method. To access data of the object, the FU

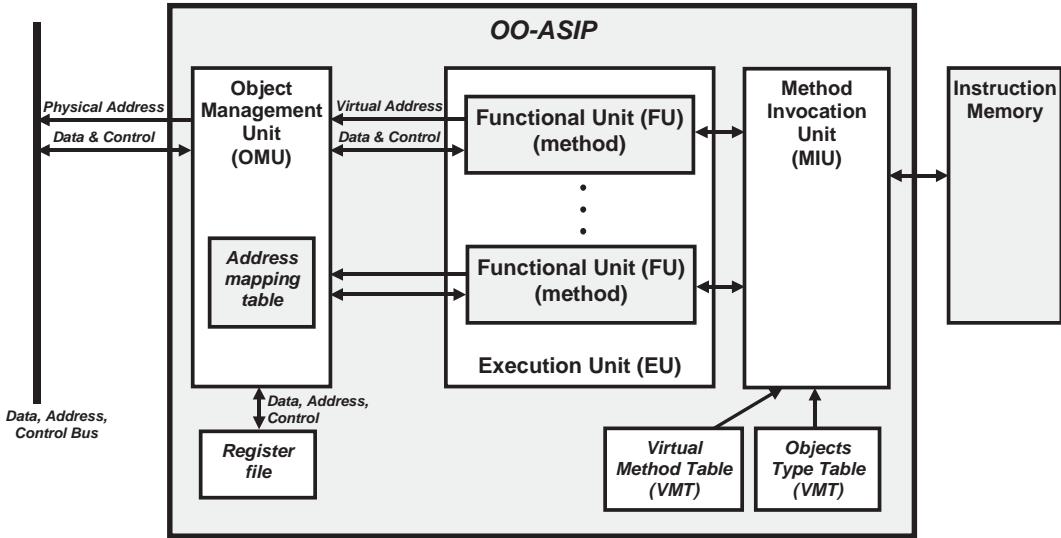


Fig. 5. Architecture of the OO-ASIP platform [45].

communicates with the Object Management Unit (OMU), which performs the memory transactions.

2.4.2. Enodia platform

Enodia [49] is an OO multi-processor platform, which uses a network structure to dispatch messages among processing elements (see Fig. 6). The Enodia platform uses firmware to dynamically resolve virtual methods to the appropriate processing elements, and provides a uniform environment for implementing a combined HW/SW OO system.

The fundamental features of the Enodia platform are as follows. Objects have a state that is stored either in the shared or private memory, and can only be changed by invoking a method on the object. For the object's state in a shared memory, all methods associated with a given object class have a uniform view of the object's state regardless of the implementation. Thus, methods may be relocated transparently between SW and HW resources, while preserving encapsulation. Methods can be physically distributed across the system; therefore, any method can be implemented in SW, HW, or FPGA. The platform uses a separate network to call methods and return the results. It also can support polymorphism by dynamically selecting the message destination for a given method call, based on the subclass of the called object.

Both analyzed platforms implement the basic concepts of OOD such as encapsulation, inheritance, polymorphism, and message passing. The differences between the OO-ASIP and Enodia platforms are as follows. (1) OO-ASIP platform implements objects as instruction-set processors, whereas Enodia platform can implement objects either in HW or in SW. (2) OO-ASIP platform represents the method invocation on an object as the instruction execution on a processor, whereas Enodia platform uses a network structure to dispatch messages among the processing elements. (3) OO-ASIP platform implements only “HW objects”, whereas Enodia

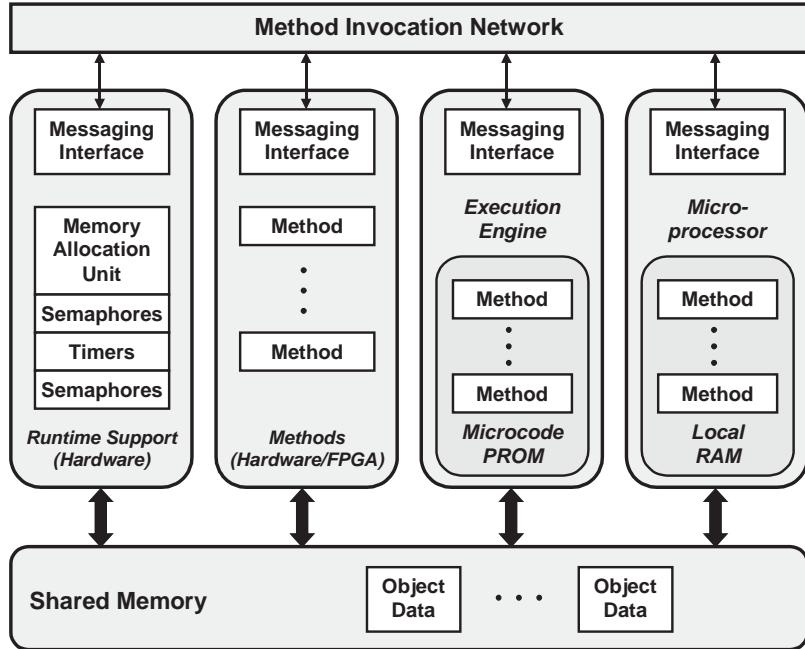


Fig. 6. Architecture of the Enodia platform [49].

platform allows HW as well as SW objects, which are treated equally and interchangeably regardless of their implementation.

2.5. HW specification using high-level formal notations

The OOD principles can also be introduced into HW domain using a formal notation of Petri Nets (PN) [50]. PN allow modeling concurrent behavior together with the data that it manipulates. In a high-level PN, the data are modeled as tokens residing on places. There have been several attempts to combine the structuring techniques of OOD with PN [51–55]. These approaches introduce classes that encapsulate token data, and functions that are executed within transitions. Currently, most existing OO PN approaches only consider the token data to be OO. However, in LOOPN++ [56], not only the token data are OO, but also the PN itself. Every element of an OO PN is an object and can be used in every context. Thus, the token data are objects that may also contain a PN, allowing very complex multi-level systems to be modeled.

The introduction of the OOD concepts into the PN notation allows developing components that encapsulate behavior and state, and supports abstraction, refinement, encapsulation, and reuse. HW components are described using classes that provide mechanisms for configuration and reuse. A system is a composition of communicating components at every level of abstraction. During system design, classes are defined, refined, reused, and configured. The abstraction and refinement by substituting less refined components with components containing an increased

amount of detail allow complex systems to be modeled and successively refined towards an implementation.

Another approach [57] aims at combining PN with UML. The authors use the PN-based behavioral specifications based on *shobi-PN v1.0* metamodel [58], instead of the UML state and activity diagrams to support the design of complex embedded systems. The metamodel presents the usual characteristics of the traditional synchronous and interpreted PN metamodel, but adds new mechanisms for supporting OO modeling and hierarchical constructs. The tokens represent objects that model the data path resources. The instance variables represent the information that is processed on the data path. The methods are the interface between the control unit and the data path. Each token models a structure of the data path. A node (a transition or a place) invokes the tokens' methods, when the tokens arrive at that node. Each arc is associated with one or more colors, which indicate the types of objects that are allowed to pass through that arc. The hierarchy can be introduced by aggregation (composition) of several objects inside one single token (a *macrotoken*) or by using the inheritance of methods and data structures.

Concurrent Object Nets (ON) [59] is a similar approach that provides a graphical semantics for classes and objects and their interaction. An ON class is a kind of graphical template for the creation of ON instances with a certain interface and behavior. The interface of an ON class consists of a set of ports. Message links connect ports of different ON instances with each other. ON instances communicate via these message links (asynchronous or synchronous) by sending simple control messages or messages with user-defined data structure.

2.6. Application of the OO techniques for HW/SW co-design

The HW/SW co-design focuses on defining both the HW and SW parts of an embedded system so that the given constraints are satisfied at minimal cost. The main topics are co-synthesis, partitioning, and co-simulation. The development and modeling of both HW and SW requires raising the level of abstraction and using a unified design methodology. The OO specifications are especially well suited for the HW/SW co-design because they provide a wealth of structuring information about the design from fine-grained methods to coarse-grained objects [12].

In the last years, co-synthesis and co-simulation from the system specifications described in the OO languages such as SystemC [15] and Java [35] has become possible. They allow producing executable specifications that can be easily and efficiently simulated. Currently, the main stream of research in HW/SW co-design focuses on the application of the OOD paradigm for *partitioning* (e.g., [60–63]).

The main advantage of having an OO system model is that it allows late binding of the components for HW/SW partitioning, because the SW models can be readily substituted with HW models and vice versa [61]. The OO specification of a system is an important aid to automatic partitioning, because it naturally describes a system at two levels of granularity. At the system-level, the system architecture is described as a network of objects that send messages between themselves to implement tasks. At the object-level, each object is described as a collection of data variables and methods that operate on the object's data. Thus, the OOD paradigm allows easy partitioning of a system into objects that could be implemented in HW or SW, as well as partitioning of objects into code and data. This gives a great opportunity for optimization and design space exploration.

The co-synthesis of the OO specifications on the *algorithmic level* is almost a standard today. But in order to raise the level of abstraction further, the partitioning at the *system-level* of abstraction is required. The OOD techniques allow implementing such a partitioning scheme [62]. It gives the designer the ability to control the partitioning process while not unnecessarily increasing the complexity. The designer can make partitioning decisions just by changing the inheritance of classes.

The HW/SW co-design promotes the cross fertilization between the HW and the SW domains, allows the unification of design efforts for system-level modeling, the application of the OOD principles to HW and embedded system design and the use of executable specifications for evaluating system requirements in its initial developments steps [63].

2.7. High-level hardware and embedded system specification using UML

UML (*Unified Modeling Language*) [16] is a standard language that is used to design, visualize, construct, model, and document SW systems. UML graphical diagrams (class, object, state, etc.) provide support for the OOD paradigm at a high level of abstraction. The application of UML for HW design is usually focused on embedded system design [17–22]. Though integrated circuits can be modeled using UML, too [64]. Recently, UML also was adopted for platform-based design [65].

An HW design can be specified in UML using three types of diagrams as follows. (1) *Class diagrams*, where behavior and data are encapsulated and structured into classes that have relationships with each other in order to implement the system functionality through interactions. (2) *Object diagrams*, where objects are instantiated from classes to implement system entities. (3) *State diagrams*, where system behavior is described in terms of an FSM. The latter, strictly speaking, is not OO.

To refine a high-level UML specification of a system (platform) into a low-level implementation (platform instance), we must define a *mapping* between the OOD concepts expressed using a subset of UML and the HDL abstractions (here we are specifically interested in VHDL). The target HDL can be either OO or not an OO one. However, in both cases, the design methodology must provide (1) a *mapping* between UML subset used to model HW and the HDL abstractions, and (2) a set of *translation rules* (and an automatic translation program, if possible) between an UML-based specification of HW and an HDL-based specification.

A mapping is usually described semi-formally using an UML *metamodel* [66] that describes the syntax of the UML diagrams using a small subset of UML. We have developed the UML–VHDL metamodel [67] consisting of a class diagram, where classes describe the syntactic components of the used UML diagram and their semantic meaning. Such metamodel provides a homomorphic mapping between UML and VHDL. Once the mapping between UML and VHDL has been defined, rules that describe the translation process between UML and VHDL can be formulated.

The aim of the translation rules is to describe how an instance of the UML metamodel (i.e., any UML model described using a subset of UML defined in the metamodel) can be transformed into an instance of a target model (i.e., a concrete VHDL specification that describes the implementation of an HW model specified using UML). Each translation-based approach usually defines its own set of rules for transformations of a metamodel. These rules can be implemented manually by an HW designer, or automatically using a dedicated translation tool.

Below, we compare two translation-based approaches: McUmber and Cheng's approach [68] focuses on HW specification using UML state diagrams, and our approach [67] focuses on the formalization of the high-level HW design processes and HW specification using UML class diagrams.

McUmber and Cheng [68] present a framework for generating VHDL specifications from a subset of UML, and a set of rules to map UML *class* and *state* diagrams to VHDL. The metamodel that describes a mapping between UML state diagrams and VHDL is briefly described below.

Each class is represented as a VHDL entity and architecture pair. Associations between classes are represented in VHDL with signals. The class instance variables are represented with the VHDL shared variables declared in the entity declarations. An UML *simple state* is represented in VHDL as a *process*. Processes representing states all have a similar structure: an initial *wait* statement waits until the *state* signal contains the name of the state. At that point, the process is activated and can perform the actions required by the state diagram.

The *events* and *messages* exchanged between objects and states are carried by signal declarations on the entity port declaration. State transitions are represented as signal assignments within processes. All events within the state diagram are represented as signals that assume a momentary value. All transitions are mapped to a loop construct that contains *wait* and *if* statements waiting for transition events.

The *composite states* are mapped to the entity–architecture pairs. The entity port description contains all the signals from its instantiating state as well as other special signals required for messages between states. The concurrent state components are mapped to concurrent composite states requiring the “thread of control” to be split between the concurrent states.

In [67], we have proposed a framework for adapting the OOD principles and, specifically, for applying the design patterns for HW design. The framework includes the generation of VHDL programs from UML class diagrams as well as a set of mapping rules that enable the generation. The simplified UML–VHDL metamodel (see Fig. 7) that describes the mapping between the UML concepts (only class diagrams are considered) and VHDL abstractions (in braces) is explained below.

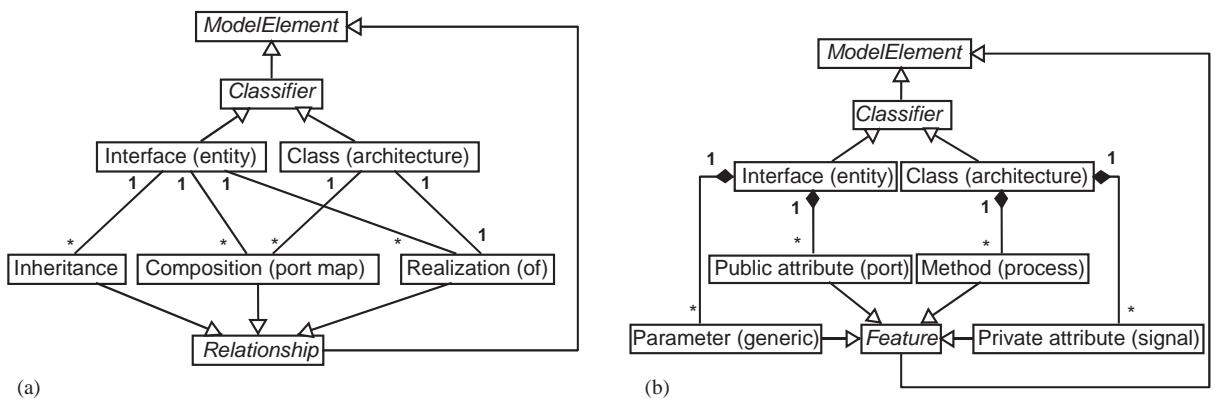


Fig. 7. UML-VHDL metamodel: (a) relationships and (b) features.

The elements of UML class diagrams are classifiers, relationships, and features. The *classifiers* are interfaces and classes that describe basic design blocks of an HW model. The *relationships* (Fig. 7(a)) describe different types of connections and associations between classifiers. The *features* (Fig. 7(b)) describe parameters, attributes and methods of classifiers.

We map an *abstract class (interface)* to a VHDL *entity*. A class that *realizes* an abstract class is mapped to the VHDL *architecture*. The class parameters are mapped to a VHDL *generic* statement, public class *attributes*—to the VHDL *ports* and private class *attributes*—to the VHDL *signals*. The class *methods* are mapped to VHDL *processes (procedures)*. The *composition* relationship describes the composition of a system from components and is mapped to a VHDL *port map* statement. The *inheritance* relationship means that an entity inherits the I/O ports from a base entity.

Both approaches have similarities as well as differences. The similarities are as follows: (1) both approaches use UML subsets as a source language and VHDL as a target language, (2) both provide mapping rules between UML and VHDL, (3) both generate the target VHDL specifications automatically.

The main difference is that our approach is oriented at reuse-based system design by adapting and integrating third-party *Intellectual Property* (IP) components, whereas the McMumber and Cheng's approach is oriented at HW design from scratch.

3. Application of design patterns for HW design

In this section, we consider the application of UML class diagrams for describing the structure of a target system at a high level of abstraction. Class diagrams can be used to describe design patterns [69]. Design patterns identify, document, and catalog successful design solutions to common design problems. Also, design patterns aid the development of reusable system models by expressing the structure and collaboration of components to designers at a level higher than source code or individual objects and classes. Their use can be extended to the development of models for HW components, embedded and real-time systems [70,71], too.

There are two approaches for adopting the concept of design pattern to HW design. The first approach aims at adapting the already known SW design patterns for HW design [72–74]. The second approach aims at discovering new HW-specific design patterns [75,76]. Below, we describe several common SW design patterns [69] and their application in HW design domain.

3.1. Union Design Pattern

The union design pattern is used to define layers of abstraction in a design. It represents the abstraction of several concrete representations of a system (component). The abstract superclass *AbstractRepresentation* represents a higher level of abstraction than the concrete subclasses. The union pattern thus defines two distinct abstraction levels. At any given moment, an operation can be performed at either the lower, more concrete abstraction layer, or at the higher, more abstract layer.

Fig. 8 presents the UML class diagram of the union design pattern. The abstract class *AbstractRepresentation* defines an abstract (i.e., not implemented, yet) method *operation1()* and

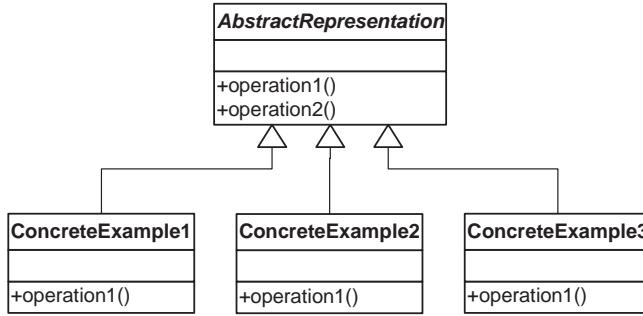


Fig. 8. Union design pattern.

implements a method *operation2()*. The *ConcreteExample* classes inherit the implementation of the method *operation2()* from *AbstractRepresentation* class and provide the implementation for the method *operation1()*.

The usage of the union design pattern allows to clearly separate between the invariant (unchanging) aspects of the design problem and the variant (changing) ones. The abstract superclass is a representation of the essence that is common to all possible concrete subclasses. It thus represents the invariant aspects of all classes in the system. The concrete subclasses are equivalent at a higher, more abstract level, and differ from each other at a lower level of abstraction. Thus, they represent the variant aspects of the design problem. The actual behavior of the system is a composition of the invariant behavior of the superclass and the variant behaviors provided polymorphically by the instances of the concrete subclasses used.

The union design pattern can be used in HW design domain to describe the configurable HW systems using, e.g., the *configuration* statement in VHDL that allows specifying a concrete implementation of a particular design entity in a system. For example, ALU can be implemented as a composition of the decoding operation (invariant aspect of a design) and abstract design entities that represent the ALU operations (variant aspects of a design). The particular implementations of the ALU operations can be specified independently, thus giving more flexibility for the designer to select the optimal implementation of ALU with respect to the required design characteristics (chip area, speed, power consumption, etc.).

3.2. State design pattern

The state design pattern is used for decoupling a system from its state and encapsulating the functionality related with this state into a separate class. The states are separated and put into different classes. The advantage is that the ability to maintain a system is much greater when behavior or algorithm, which is performed when a system enters a certain state, is encapsulated into its own class, rather than hidden somewhere in the body of some method.

Fig. 9 presents the UML class diagram of the state design pattern. The class *Context* encapsulates a state. The abstract class *AState* describes a representation of the state's functionality. The *ConcreteState* classes provide an implementation of the particular states.

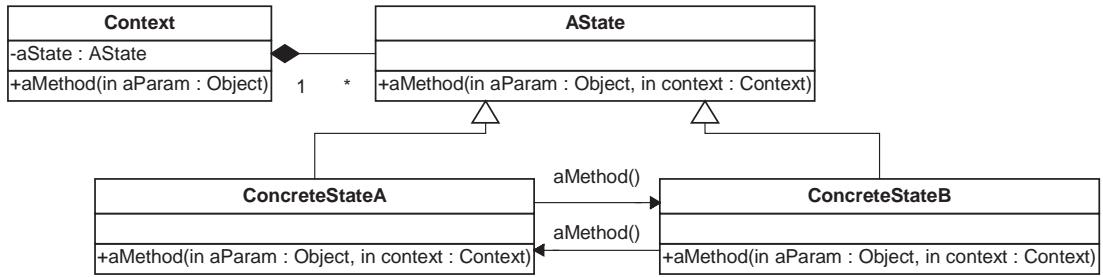


Fig. 9. State design pattern.

The state design pattern can be used in HW domain to design FSMs. Each *ConcreteState* class provides an implementation of a separate FSM state and the corresponding functionality that is performed when the FSM enters this state. This solution allows to design FSMs that can be easily reused and extended with new states for other future applications.

A disadvantage of the state design pattern is that it may result in excessive complexity of the designed system since for every state a separate object is assigned and later implemented as a separate component. Thus, the application of this pattern leads to worse performance characteristics of the designed system. The trade-off here is between faster implementation of the system and worse characteristics.

3.3. Diamond design pattern

The diamond design pattern is used to represent the behavioral complexity of a design that provides multiple services for its environment. To achieve composition of service classes, multiple inheritance is used. It allows a derived class to obtain features from two or more parent classes.

Fig. 10 presents the class diagram of the diamond design pattern. The abstract class *aService* describes an interface that is common for all services provided by the system. The *cService* classes inherit from *aService* and provide an implementation of the concrete services. The class *System* inherits the implementation of services from the *cService* classes and describes the multi-service system.

The diamond pattern can be widely used in HW and embedded system design domains, e.g., for composing new systems from the available HW and/or SW components. The disadvantage of the pattern is that it may introduce considerable performance overhead, because not all services of the classes composed using inheritance may be required in a designed system.

3.4. Wrapper/decorator design pattern

The wrapper (or sometimes called *Decorator*) design pattern is used for adding additional functionality to a particular component. A wrapper (or decorator) is a component that has some parts of an interface identical to the interface of the contained component. Any calls that the wrapper gets, it sends to the component that it contains, and adds its own functionality, either before or after the call. This allows a lot of flexibility and reusability, since the designer can change

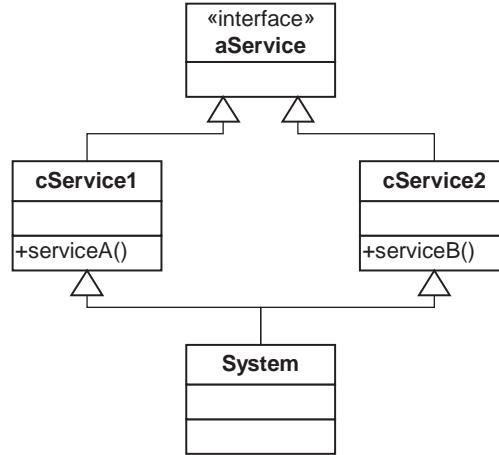


Fig. 10. Diamond design pattern.

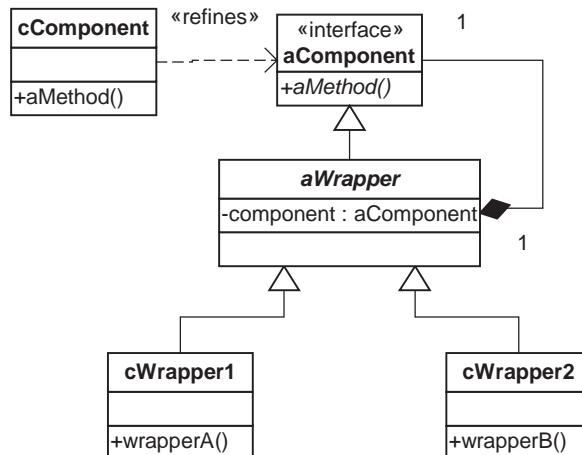


Fig. 11. Wrapper design pattern.

the wrapper without changing the original component it wraps. As the designer can recursively nest wrappers, the pattern allows for almost infinite amount of customization.

Fig. 11 presents the UML class diagram of the wrapper design pattern. The abstract class *aComponent* specifies an interface that is common for all components and their wrappers. The class *cComponent* provides an implementation for *aComponent* class. The abstract class *aWrapper* specifies a wrapper interface and contains an instance of *aComponent* class. The *cWrapper* classes provide the different implementations of a wrapper.

The wrapper pattern can be widely used in HW design, especially for (1) communication control to provide an implementation of a communication protocol for an HW component communicating with its environment, (2) fault-tolerant applications to implement fault-avoidance

and fault-discovery techniques such as TMR (*Triple-Modular Redundancy*) and BIST (*Built-In Self Test*), (3) adaptation of the physical memory interfaces to a communication network that may have a different number of access ports, and (4) layered data packet processing in network applications.

A disadvantage of the wrapper design pattern is that it may result in considerable overhead since the complete interface of the wrapped component needs to be handled by the wrapper, including those interface elements that need not be adapted. Thus, the application of the wrapper design pattern may lead to excessive amounts of the adaptation code and performance overhead.

3.5. Composite design pattern

The composite pattern allows treating both single components and collections of components identically. It allows building complex components by recursively composing similar components in a hierarchical tree-like manner. It also allows the components to be manipulated in a consistent manner, by requiring all of the components to have a common parent. The composite design pattern is often used to represent recursive system architectures and their behavior.

Fig. 12 presents the UML class diagram of the composite design pattern. The abstract class *aComponent* specifies an interface that is common to all its implementations. The *Leaf* classes provide the different implementations of *aComponent*. The class *Composite* is recursively composed of an array of *aComponent* type components (which may be of *Leaf* or *Composite* type) and implements the recursive call to their operations.

The composite design pattern can be widely used in HW design to rapidly implement systems with a repetitive architecture such as multiplexer arrays, register arrays, adders, etc. The advantage of using the composite design pattern is scaling, i.e., a part of the system can be addressed in the same way as the whole system. The disadvantage of this pattern is, as it is with many other design patterns, the possibility of the considerable performance overhead introduced by the recursive composition of the components.

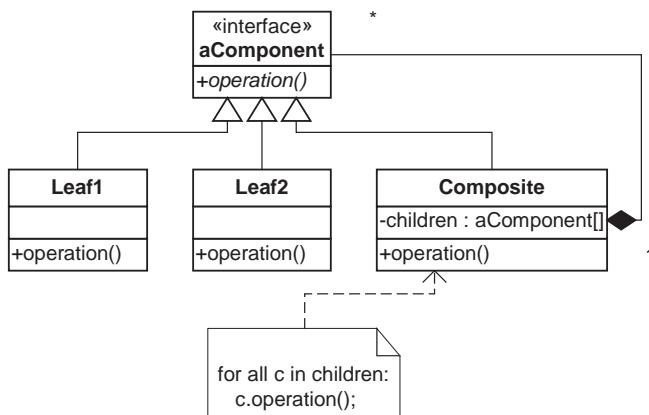


Fig. 12. Composite design pattern.

4. Principles of OOD of HW and embedded systems

In this section, we summarize our overview and formulate the basic principles of the OOD as applied to HW and embedded system design domains.

Principle 1. OOD is based on representation of the real-world entities using a set of classes and relationships between them.

Principle 1 states that domain entities are represented as OO models. In order to represent the domain and its systems, a higher layer of abstraction is introduced above the domain abstraction layer. The main actors of this higher abstraction layer are classes—the highly abstract concepts that encapsulate the domain data and operations performed on it. The classes are related with each other through a network of relationships. These relationships represent different types of classification and interaction between domain entities—the abstract ones introduced during domain analysis as well as the concrete ones that represent direct communication between domain entities. The decomposition of a domain system into classes and relationships between them is a matter of an OO domain analysis.

Principle 2. OOD is an evolutionary design methodology specifically oriented at design for reuse.

The main concepts of the OOD are essentially the same as in other design methodologies such as component-based design. What is the difference is the way these concepts are applied. The OOD is essentially the evolutionary design methodology specifically aimed for implementing the ‘design-for-reuse’ principle systematically, whereas other design methodologies usually support the design reuse only casually and ad hoc. The classes are specifically designed to evolve over time as the designer analyzes and better understands the domain, as well as new requirements appear in the process of system design. The classes can be extended using a powerful extension mechanism of *inheritance* that allows adding new data and methods to the existing classes. As the designer adds new features to a system, he/she may discover that his/her previous design does not support easy system extension. With this new information, the designer can restructure parts of the designed system, very possibly adding new classes or class hierarchies to an OO model of a system.

Principle 3. The OO model of a system is an instance of the class hierarchy that abstracts the domain and consists of the communicating objects.

The class hierarchy organizes the commonalities present in a set of domain entities into a tree structure, where the root of any subtree contains all the attributes and behavior common to every descendant of that root. In standard HDLs, there are no syntactical structures to express this kind of hierarchy, therefore reuse is limited to the component level, i.e. designers must use the component “as is” or have to design a new one. In the OOD paradigm, *inheritance* provides the potential for building new solutions to problems by adding incremental capability to existing problem solutions through subclassing. This concept is one of the most fundamental OOD features to enable design reuse in the domain.

An object is a distinct instance of a given class that encapsulates its implementation details and is structurally identical to all other instances of that class. Multiple instantiations of a given class can be made and each of them represents a distinct object. An object defines an interface of services (operations, methods) that it provides, while encapsulating the data defined in its class

and hiding the implementation details. In HW design domain, the notion of object corresponds to a distinct component instance that is used to create unique references to the lower-level components.

Principle 4. The functionality of the OO system model is defined by the relationships between objects (classes).

There are three basic types of relationships that describe the structural affinity between classes and behavioral interaction between objects in the OO model:

- (1) *Inheritance* allows a subclass to incorporate the data and operations defined in a parent class and to augment it with the additional data elements and methods. In HW domain, it corresponds to the functional classification of components.
- (2) *Aggregation* is used when one object (container) physically or conceptually contains another object (component). In HW design domain, it corresponds to the composition of components into a larger module (system). The purpose of the component is to contribute in some way to the overall functionality of its container (system).
- (3) *Association* represents conceptual relationships between classes, and it is used for exchange of messages. In HW domain, it corresponds to the dependence of a particular component upon another component for interaction or communication.

Principle 5. OO models of the domain systems can have three different implementation views: structural, behavioral, and FSM-based.

Depending upon the implementation on the domain layer of the abstraction, the OO HW systems can have three distinct implementation views as follows:

- (1) The *structural* view to the HW system is based on the representation of *physical components* as objects. It can be represented using UML *class diagrams*.
- (2) The *behavioral* view to the HW system is based on the representation of *abstract concepts* for solving the design problem. The objects of the design process are abstract *design processes*. It can be represented using UML *object diagrams*.
- (3) The *FSM-based* view to the HW system represents each object as an FSM. It can be represented using UML *state diagrams*.

Principle 6. The common OO HW models can be expressed in terms of the OOD paradigm using class templates and polymorphism.

In UML, a *parameterized class (template)* is used to express genericity in the OO models. Templates allow modeling generic domain components explicitly and thus providing an alternative approach for representing generalization in the domain: not only via class hierarchies, but also using class templates and their instances.

Polymorphism is another mechanism provided by the OOD methodology that provides the ability to manipulate with objects of distinct classes using only knowledge of their common properties without regard for their exact class type. The polymorphism enables new classes possessing required properties to be handled equally with their base class. Polymorphism, when applied to HW design, gives an enormous potential of reuse enabling the incremental refinement of particular HW modules.

Principle 7. Common domain architectures can be expressed in terms of the OOD paradigm using design patterns.

Common OO solutions to the HW domain problems can be expressed more abstractly and generally using *HW design patterns*. The difference of the design patterns from the other OO models and class hierarchies is that a design pattern is a more general and abstract model that represents a common and well-proven solution to a design problem rather than a specific domain system. Thus, design patterns represent an abstraction layer that is *higher* than common OO domain models, and they allow fostering reuse of design knowledge and experience rather than the reuse of a specific implementation.

Principle 8. An object/domain metamodel is the conceptual framework for connecting higher and lower layers of the abstraction in system design.

The OO domain models specify only abstract views of domain systems without specifying the particular domain content. This low-level content must be introduced later in the design process when OO models are refined into domain entities.

Thus, the introduction of the *higher-level model* that abstracts the domain content on the lower layer of abstraction, requires the introduction of the *object/domain metamodel* that describes mapping between those layers of abstraction within the OOD framework. The OO layer of the domain system can be described using a graphical notation such as UML diagrams or an OO language. The transition to the domain layer can be performed using the traditional compilation techniques or generation into an HDL using the generative technology. All these different choices greatly influence the design productivity as well as the efficiency of the designed HW system.

5. Case studies

We illustrate the application of the OOD principles discussed above by presenting two simple case studies. The first one implements a diamond pattern for designing multi-protocol communication interface for HW components described in VHDL. The second one implements a composite pattern for design of Network-on-Chip model described in SystemC. With these two case studies, we aim to demonstrate the usage of two most popular domain languages, VHDL that is not an OO HDL and SystemC that is an OO HDL, in the context of the OOD of HW components.

5.1. Specification of communication control using Diamond pattern

In this case study, we use a diamond design pattern (see Section 3.3) to design an HW component with a multi-protocol communication service provided for communication control. The services are composed of an HW component, which implements basic functionality, using a wrapper design pattern [74].

The main purpose of communication control is to ensure the relevant transmission of data (e.g., operands, commands, addresses, etc.) to and from the HW component. The transmission of data can be described using different rules or *protocols*, i.e. an agreed format for transmitting data between the HW components. We consider here two common communication protocols, namely,

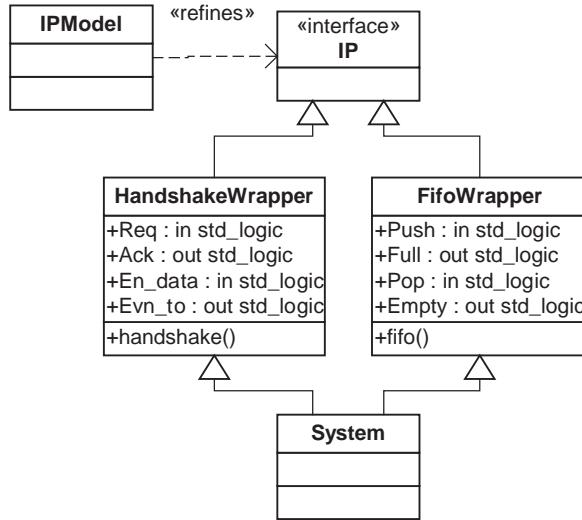


Fig. 13. Class diagram of a target system.

handshake protocol that deals with an asynchronous flow of data, and FIFO protocol that deals with sudden bursts of data in a producer–consumer model. Our aim is to include both into a target system and let the user decide which protocol should be used to communicate with the wrapped HW component.

Fig. 13 presents the UML class diagram of a target system. Abstract class *IP* is an interface of an HW component (soft IP). Class *HandshakeWrapper* provides an implementation of a handshake protocol and class *FIFOWrapper* provides an implementation of a FIFO protocol. Class *System* uses multiple inheritance to inherit both protocol implementations and encapsulates an instance of *IP* component.

The architecture of the target system is shown in Fig. 14. The architecture implements the FIFO-based model of an object described in Section 2.4. Note that signals *Data_in* and *Data_out* represent the IP-specific data signals, and signal *Method* is used to select the required method of the object.

IP is an HW component (e.g., ALU, controller, processor, etc.) for which a multi-protocol communication interface must be implemented. The handshake wrapper wraps the *IP* with handshake FSM. The data are transferred to the *IP* when request signal *Req* is set to a high level and acknowledgement signal *Ack* is received from the handshake FSM. The results are returned as soon as the *IP* processes the data.

The FIFO wrapper wraps the *IP* component with two instances of FIFO buffer and additional control logic. The internal clock signal *clk_int* is used to run FIFO control logic and the IP. The data are transferred to the IP component when *Push* signal is set to a high level and *Full* signal has a low level. The results are returned when *Pop* signal is set to a high level and *Empty* signal has a low level.

We have implemented the target system using several freely available HW components described in VHDL as the IP. The synthesis results (Synopsys; CMOS 0.35μm technology) are

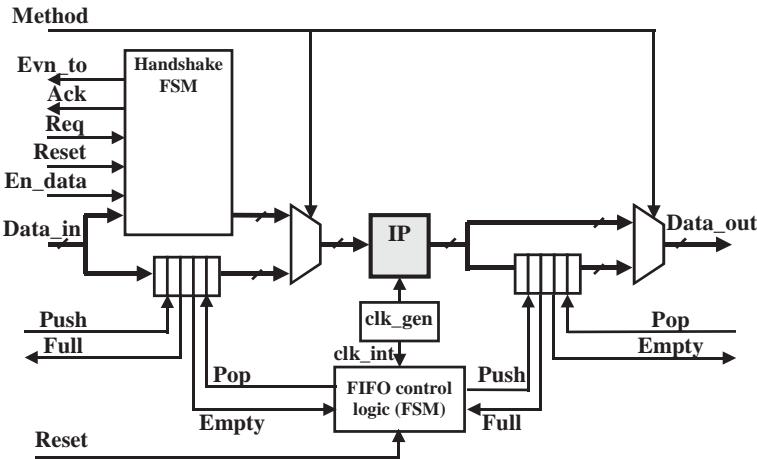


Fig. 14. Architecture of a target system.

Table 1
Synthesis results

Soft IP	Area, cells (soft IP)	Area, cells (system)	Increase (%)
Free-6502 [77]	4670	7423	59
Dragonfly [78]	5883	11832	101
HC11 [79]	7394	15494	110
AX8 [80]	8020	13219	65
i8051 [81]	24258	30529	26

presented in Table 1. The results show that chip area of the target system is on average about 72% larger than the original HW component.

5.2. Specification of a Network-on-Chip using composite pattern

In this case study, we apply the composite design pattern (see Section 3.5) to model a *Network-on-Chip* (NoC) [82]. NoC is a novel communication-oriented SoC design platform based on the application of network technologies for connecting HW components on a chip. A system is designed as a network of nodes that communicate one with another by exchanging data packets. Each packet carries the encapsulated data, source node ID, target node ID, and other

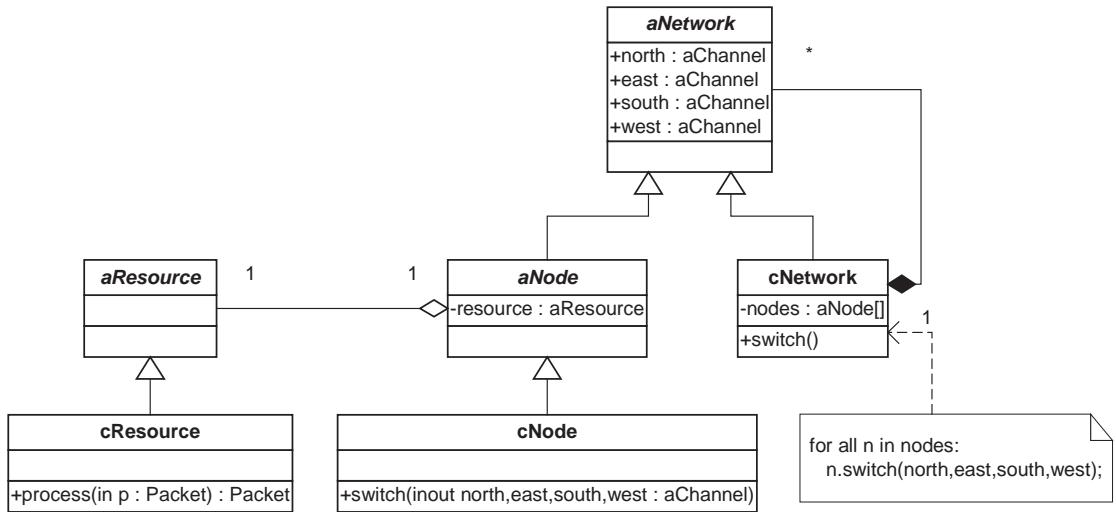


Fig. 15. Class diagram of an NOC.

network-specific information. Each node contains a resource that represents a particular component of a system. Sending data packets to a particular network node and receiving the results perform the necessary computations required for the specific application of NoC.

Fig. 15 presents the UML class diagram of the NoC. The abstract class *aNetwork* describes a part of NoC that has four channels to communicate with other parts of an NoC. Class *aChannel* is an abstraction of the communication protocol. A derived class *aNode* inherits channels from *aNetwork* class and contains an instance of an abstract class *aResource*, which describes a resource provided by an NoC node. Class *cResource* is a concrete resource of a concrete node *cNode*. Method *switch(...)* performs data packet switching until a packet reaches its destination. Then, method *process()* processes the packet, performs the computations on the data that is carried by the received packet, encapsulates the results in a new packet, and forwards it back to the sender of the original data packet. Class *cNetwork* is a concrete NoC that consists of an array of nodes connected with channels that are switched simultaneously.

The NoC architecture (see Fig. 16) provides the communication infrastructure for the resources on a chip. An NoC consists of *resources* and *switches* that are connected using *channels* in mesh network architecture. The nodes can communicate with each other by sending messages. A resource is a computation or storage unit or their combination, e.g., embedded processor, DSP core, RAM, or application-specific HW resources. Switches are used to route and buffer messages between resources. Each switch is connected to four other neighboring switches through channels. A channel consists of two one-directional peer-to-peer buses between two switches.

Fig. 17 presents the example of an NoC node specification in SystemC. Fig. 17(a) shows the structure of a data packet and Fig. 17(b) presents an NoC node that inherits from class *aNode* and implements the data packet switching. *Sc_fifo<Packet>* is a channel that uses FIFO to buffer the transmitted packets.

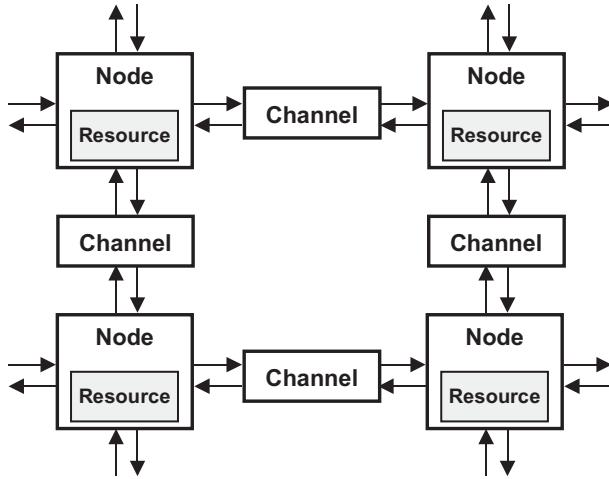


Fig. 16. Typical architecture of an NOC.

<pre> struct Packet { sc_bit<4> sourceID; sc_bit<4> targetID; sc_bit<8> op_code; sc_bit<16> data1; sc_bit<16> data2; sc_bit<4> packetNo; sc_bit<1> checksum; sc_bit<1> error; sc_bit<1> fill; }; </pre> <p style="text-align: center;">(a)</p>	<pre> template <class Packet> class cNode: public aNode<Packet> { public: void switch(sc_fifo<Packet> north, sc_fifo<Packet> east, sc_fifo<Packet> south, sc_fifo<Packet> west); SC_CTOR(cNode) { SC_METHOD(switch); } }; </pre> <p style="text-align: center;">(b)</p>
--	--

Fig. 17. Specification (a fragment) of NoC in SystemC: (a) data packet and (b) network node.

6. Evaluation and discussion

Current SoC and embedded system design methodologies are not sufficient to cope with more and more demanding requirements to system design such as design reuse and time to market. The OOD methodology is a strong candidate to resolve these problems. We have identified the following approaches for adopting the OOD paradigm for HW design. Below, we summarize their advantages and disadvantages.

- (1) The extension of an existing HDL with the abstractions that implement the OOD concepts. The advantage is its strong HW foundation. The disadvantage is the lack of tools to support debugging, verification, and synthesis for these languages.
- (2) The usage of an existing language that supports the OOD paradigm and adds new constructs that support concurrent behavior and other needed HW modeling constructs. The advantage is the already available support for the OOD concepts. The disadvantage of this approach is

that this addition of constructs to the language requires an extension or even a redefinition of the language, and the development of new tools for simulation.

- (3) The implementation of the OO HW architectures allows selecting several different implementation models such as behavioral, structural and FSM-based. The advantage is that a designer can implement exactly those OOD concepts, which he needs. The disadvantage is that the design of OO HW architectures at a low-level of abstraction may be tedious and error-prone.
- (4) The usage of UML for high-level specification of HW. The translation from the OO model to the non-OO implementation is performed during the generation of an HW model. The major advantage of this approach is the usage of existing tools to support debugging, verification, and synthesis for these languages. Other advantages of using UML for HW design are as follows: (a) high-level specification of a designed system, (b) better soft IP reusability and adaptability, and (c) better documentation for further reuse and maintenance of a system. The major disadvantage is the complexity of the transformation of an OO model and difficult verification problems. Other difficulties of using UML for HW design are related with (a) specification of interconnections between components, (b) specification of generic domain functionality, and (c) increased initial development time.

We summarize the current achievements in the area of the application of the OOD concepts in HW and embedded system domain as follows.

- (1) The introduction of SystemC language spurred the interest of HW designers and researchers to adopt more and more of the OOD concepts in SoC design domain. There is a continuous effort to develop new versions of SystemC by its authors as well as to provide extensions to SystemC by other researchers that widen a set of available OOD features for HW modeling and synthesis.
- (2) There is a continuous effort to further develop and extend VHDL in order to adapt it to the new level of abstractions emerging in HW design, including the object level of abstractions. A plethora of extensions to VHDL has emerged and continues to be developed, which ensures that this powerful domain-specific language continues to remain a strong competitor to SystemC in the future.
- (3) The research in the OO HW architectures and platforms has been based on further exploiting the well-known FSM-based model for representing objects in RTL. An FSM is a flexible model of computation that is known very well by HW designers and can be easily implemented by synthesis tools.
- (4) The designers are starting to use more widely the graphical OOD diagrams described using a standard UML language for HW and embedded system design, in addition to the casual block-based diagrams. New tools are developed that allow generating the corresponding HDL code from the UML diagrams automatically.

The advantages of using the OOD concepts in HW design are as follows:

- (1) *Classes and design patterns* introduced into HW design from the SW engineering domain are common for designing SW/HW parts of an electronic system. Thus they can contribute to the unification of design methodologies and affect the merging of SW/HW design at a higher level of abstraction.

- (2) *Design patterns* allow for specifying an HW design problem at a higher abstraction level graphically. The application of design patterns described using UML class diagrams allows for specifying target systems using abstract representations of proven design experience. The design content is captured immediately and intuitively, thus increasing design comprehensibility. The pattern-based design can be easily supported by the automated validation and code generation tools, thus increasing design reuse, quality, and productivity. The level of abstraction is raised to the system level, which allows dealing with growing complexity of HW designs.
- (3) *Encapsulation* allows to separate the data from functionality, thus implementing the separation of concerns principle in HW design.
- (4) *Inheritance* and *polymorphism* are powerful OOD principles that allow to increase design productivity substantially, because they enable reuse of existing code.
- (5) *Message exchange* is a very promising OOD concept especially when considered in the context of a new emerging paradigm for the interconnection-based SoC design, called NoC [82].

The idea of adopting the OOD principles to HW and embedded system design domain holds great promises; however, there are many gaps that must be bridged before these promises could be fully fulfilled and they are as follows:

- (1) *Conceptual gap*: SW designers think in terms of the OOD concepts (objects and messages), whereas HW designers are used to think in terms of the component-based design (components and wires).
- (2) *Physical gap*: how the physical constraints (such as the timing ones) should be reflected in an OO model or a metamodel (design pattern)?
- (3) *Technological gap*: how objects or even entire design patterns could be directly synthesized to RTL?

7. Conclusions and future research

Currently, the design and implementation of HW, embedded and real-time systems hardly differ from traditional SW development. Despite all known problems and difficulties, the OOD principles can be adopted for developing HW as well as embedded systems. The OOD paradigm provides the techniques like abstraction, inheritance and polymorphism that help the designer to handle large and complex systems. On the other hand, the OOD paradigm does not provide the demanded code efficiency (in terms of chip area or delay) for real-time and embedded applications.

The main task of the future research on OO HW design, therefore, should focus on reconciling these two contradicting features of the OOD paradigm, and seeking for the best possible trade-off solutions between faster design time provided by the application of the OOD paradigm and worse system characteristics introduced by the OO abstractions. We formulate the most probable future trends in the application of the OOD principles to HW and embedded system design as follows.

- (1) A great deal of research is performed to extend UML in order to take new domains of application into account [21]. The research focuses on the UML extensions for modeling (a) real time behavior, (b) executable specifications capable to express parallelism, (c) HW/SW

interfaces and co-simulation, (d) behavior of analog interfaces and systems, and (e) size, weight, area, power, timing and other constraints.

- (2) The adoption of design patterns for HW design will attract more and more attention from the researchers. The efforts will be directed at adopting the already known SW design patterns as well as at discovering new HW-specific design patterns.
- (3) The researchers are seeking for the unification of HW and SW design methodologies into a single SoC design flow under the OOD paradigm umbrella.

References

- [1] H. Chang, L. Cooke, M. Hunt, G. Martin, A. McNelly, L. Todd, *Surviving the SOC Revolution: A Guide to Platform-Based Design*, Kluwer Academic Publishers, Dordrecht, 1999.
- [2] SEMATECH, *The International Technology Roadmap for Semiconductors*, 2001.
- [3] F. Vahid, T. Givargis, *Embedded System Design: A Unified Hardware/Software Introduction*, Wiley, New York, 2002.
- [4] A. Pawlak, W. Włodzimierz, Modern object-oriented programming language as a HDL, in: Eighth IFIP International Symposium on Computer Hardware Description Languages and their Applications, 1987, Amsterdam, pp. 343–362.
- [5] W.H. Wolf, How to build a hardware description and measurement system on an object-oriented programming language, *IEEE Trans. Comput. Aided Des.* 8 (3) (1989) 288–301.
- [6] R. Gupta, W.H. Cheng, R. Gupta, I. Hardonag, M.A. Breuer, An object-oriented VLSI CAD framework: a case study in rapid prototyping, *IEEE Comput.* 22 (9) (1989) 28–37.
- [7] A.C. Verschueren, An object-oriented design and simulation system for VLSI”, *Microprocess. Microprogramming* 30 (1–5) (1990) 241–246.
- [8] D.C. Craig, Circuit description and elementary hierarchical circuit simulation using C++ and the object-oriented programming paradigm, Ph.D. Dissertation, Department of Computer Science, Memorial University of Newfoundland, St. John’s, Newfoundland, Canada, April 1991.
- [9] S. Kumar, J.H. Aylor, B.W. Johnson, W.A. Wulf, Object-oriented techniques in hardware design, *IEEE Comput.* 27 (6) (1994) 64–70.
- [10] K. Agsteiner, D. Monjau, S. Schulze, Object-oriented system level specification and synthesis of heterogeneous systems, *Proceedings of EURO-MICRO’95*, Como, Italy, 4–7 September 1995.
- [11] W. Nebel, G. Schumacher, Object-Oriented Hardware Modelling—Where to apply and what are the objects? *Proceedings of EURO-DAC with EURO-VHDL’96*, Geneva, Switzerland, September 16–20, 1996, pp. 428–433.
- [12] W.H. Wolf, Object-oriented cosynthesis of distributed embedded systems, *ACM Trans. Des. Automati. Electron. Systems* 1 (3) (1996) 301–314.
- [13] J.M. Bergé, O. Levia (Eds.), *Object-Oriented Modeling Current Issues in Electronic Modeling*, Vol. 7, Kluwer Academic Publishers, Dordrecht, 1996.
- [14] T. Potok, M. Vouk, Development productivity for commercial software using object-oriented methods, *Proceedings of the 1995 Conference of the Centre for Advanced Studies on Collaborative Research*, IBM Press, Toronto, Ontario, Canada, 1995, p. 52.
- [15] S. Swan, D. Vermeersch, P. Hardee, T. Hasegawa, A. Rose, A. Coppola, M. Jansen, T. Grötker, A. Ghosh, K. Kranen, Functional Specification for SystemC 2.0, White paper, OSCI, 2001.
- [16] G. Booch, I. Jacobson, J. Rumbaugh, J. Rumbaugh, *The Unified Modeling Language User Guide*, Addison-Wesley, Reading, MA, 1998.
- [17] J.M. Fernandes, R.J. Machado, H.D. Santos, Modeling Industrial Embedded Systems with UML, *Proceedings of the Eighth IEEE/IFIP/ACM International Workshop on Hardware/Software Co-Design (CODES’2000)*, San Diego, CA, USA, May, 2000, pp. 18–23.
- [18] G. de Jong, A UML-based design methodology for real-time and embedded systems, *Proceedings of the Design Automation and Test in Europe (DATE 2002)*, Paris, France, 4–8 March 2002, pp. 776–778.

- [19] Q. Zhu, A. Matsuda, S. Kuwamura, T. Nakata, M. Shoji, An object-oriented design process for system-on-chip using UML, Proceedings of the 15th International Symposium on System Synthesis (ISSS 2002), Kyoto, Japan, 2002, pp. 249–254.
- [20] G. Martin, UML for embedded systems specification and design: motivation and overview, Proceedings of DATE 2002, Paris, France, 4–8 March 2002, pp. 773–775.
- [21] S. Gerard, F. Terrier, UML for real-time, in: L. Lavagno, G. Martin, B. Selic (Eds.), UML for Real, Kluwer Academic Publishers, Boston, 2003, pp. 17–52.
- [22] M. Edwards, P. Green, UML for hardware and software object modeling, in: L. Lavagno, G. Martin, B. Selic (Eds.), UML for Real, Kluwer Academic Publishers, Dordrecht, 2003, pp. 127–148.
- [23] IEEE. Standard VHDL Language Reference Manual, IEEE Std 1076 1993, 1994.
- [24] W. Ecker, An object-oriented view of structural VHDL description, Proceedings of VHDL International Users' Forum (VIUF Spring '96), Santa Clara, CA, USA, March 1996, pp. 255–264.
- [25] S. Swamy, A. Molin, B. Covnot, OO–VHDL: object-oriented extensions to VHDL, IEEE Comput 28 (10) (1995) 18–26.
- [26] R. Zippelius, K.D. Muller-Glaser, An object-oriented extension of VHDL, Proceedings of the VHDL-Forum Spring'92 Meeting, Santander, Spain, April 26–28, 1992, pp. 155–163.
- [27] W. Glunz, A. Pyttel, G. Venzl, System-level synthesis, in: P. Michael, U. Lauther, P. Duzy (Eds.), The Synthesis Approach to Digital System Design, Kluwer Academic Publishers, Dordrecht, 1992, pp. 221–260.
- [28] J. Benzakki, B. Djafri, Object oriented extensions to VHDL, the LaMI proposal, Proceedings of the International Conference on Hardware Description Languages and their Applications, Toledo, Spain, 1997, pp. 334–347.
- [29] M. Radetzki, W. Putzke-Röming, W. Nebel, Objective VHDL: The object-oriented approach to hardware reuse, in: J.-Y. Roger, B. Stanford-Smith, P.T. Kidd (Eds.), Advances in Information Technologies: The Business Challenge, IOS Press, Amsterdam, 1998, pp. 796–803.
- [30] P.J. Ashenden, Overview of SUAVE language features, in: Proceedings of the Forum on Design Languages (FDL), Lyon, France, August 30–September 3, 1999, pp. 269–278.
- [31] G. Schumacher, W. Nebel, Survey on object-oriented languages for hardware design methodologies, in: High-Level System Modeling: Specification Languages, Current Issues in Electronic Modeling, Vol. 3, Kluwer Academic Publishers, Dordrecht, 1995.
- [32] G. Schumacher, Object-oriented hardware specification and design with a language extension to VHDL, Ph.D. Thesis, Universität Oldenburg, Germany, 1999.
- [33] M. Radetzki, Synthesis of digital circuits from object-oriented specifications, Ph.D. Thesis, Universität Oldenburg, Germany, 2000.
- [34] S. Sutherland, SystemVerilog For Design: A Guide to Using SystemVerilog for Hardware Design and Modeling, Kluwer Academic Publishers, Norwell, 2004.
- [35] T. Kuhn, W. Rosenstiel, U. Kebschull, Description and simulation of hardware/software systems with Java, Proceedings of the 36th Design Automation Conference (DAC'99), New Orleans, LA, USA, 21–25 June, 1999, pp. 790–793.
- [36] R. Allen, D. Gajski, The case for C/C++ hardware design, EETimes, 2000.
- [37] D. Verkest, J. Kunkel, F. Schirrmeyer, System level design using C++, Proceedings of the Design, Automation and Test in Europe (DATE 2000), Paris, France, 27–30 March 2000, pp. 74–81.
- [38] D. Ku, G. De Micheli, HardwareC: a language for hardware design, Technical Report SCSL/CSL/TR-90-419, Computer Systems Laboratory, Stanford University, Stanford, CA, USA, August 1990.
- [39] J. L. da Silva, C. Ykman-Couvreur, G. de Jong, Matisse: a concurrent and object-oriented system specification language, Proceedings of the IFIP International Conference on Very Large Scale Integration (VLSI'97), Gramado, Brazil, August 26–29, 1997.
- [40] A. Gerstlauer, R. Domer, J. Peng, D. Gajski, System Design—A Practical Guide with SpecC, Kluwer Academic Publishers, Dordrecht, May 2001.
- [41] T. Grötker, S. Liao, G. Martin, S. Swan, System Design with SystemC, Kluwer Academic Publishers, Boston, 2002.
- [42] W. Müller, W. Rosenstiel, J. Ruf, SystemC: Methodologies and Applications, Kluwer Academic Publishers, Dordrecht, 2003.

- [43] A. Tsikhanovich, E.M. Aboulhamid, Object-oriented techniques in hardware modeling using SystemC, in: Northeast Workshop on Circuits and Systems, Montréal, Canada, June 2003.
- [44] S. Virtanen, D. Truscan, J. Lilius, SystemC based object oriented system design, in: Proceedings of the fourth International Forum on Design Languages (FDL'01), Lyon, France, September 3–7, 2001.
- [45] C. Schulz-Key, M. Winterholer, T. Schweizer, T. Kuhn, W. Rosenstiel, Object-oriented modeling and synthesis of SystemC specifications, in: Proceedings of the Asia South Pacific Design Automation Conference (ASPDAC'04), Yokohama, Japan, 2004, pp. 238–243.
- [46] E. Grimpe, F. Oppenheimer, Aspects of object oriented hardware modelling with SystemC-plus, in: A. Mignotte, E. Villar, L. Horobin (Eds.), System on Chip Design Languages, Extended Papers: Best of FDL'01 and HDLCon'01, Kluwer Academic Publishers, Dordrecht, 2002, pp. 213–223.
- [47] A. Sangiovanni-Vincentelli, G. Martin, A vision for embedded systems: platform-based design and software methodology, *IEEE Des. Test Comput.* 18 (6) (2001) 23–33.
- [48] M. Goudarzi, S. Hessabi, A. Mycroft, Object-oriented ASIP design and synthesis, in: Forum on Specification and Design Languages (FDL'03), Frankfurt, Germany, September 23–26, 2003.
- [49] Silicon Infusion Ltd., An Object Oriented Re-Configurable System on Chip, White paper, 2003.
- [50] J.L. Peterson, Petri Net Theory and the Modeling of Systems, Prentice-Hall, Englewood cliffs, NJ, 1981.
- [51] R. Esser, An object-oriented petri net language for embedded system design, Proceedings of the Eighth International Workshop on Software Technology and Engineering Practice (STEP'97), London, UK, July 1997, pp. 216–223.
- [52] C. Lakos, The object orientation in object petri nets, in: Proceedings of the 16th International Conference on Application and Theory of Petri Nets, First Workshop on Object-Oriented Programming and Models of Concurrency, Torino, Italy, 1995 pp. 1–14.
- [53] P.A.C. Verkoulen, Integrated information systems design: an approach based on object oriented concepts and petri nets, Ph.D. Thesis, University of Technology, Eindhoven, the Netherlands, 1993.
- [54] Y.K. Lee, S.J. Park, OPNETS: an object-oriented high-level Petri-net model for real-time system modeling, *J. Systems Software* 1 (20) (1993) 69–86.
- [55] R.J. Machado, J.M. Fernandes, A.J. Proença, Specification of industrial digital controllers with object-oriented petri nets, *IEEE International Symposium on Industrial Electronics (ISIE'97)*, Vol. I, Guimarães, Portugal, July, 1997, pp. 78–83.
- [56] C.A. Lakos, C.D. Keen, LOOPN++: a new language for object-oriented petri nets, *Proceedings of the Modelling and Simulation (European Simulation Multiconference)*, Barcelona, Spain, 1994, pp. 369–374.
- [57] R.J. Machado, J.M. Fernandes, H.D. Santos, a methodology for complex embedded systems design: petri nets within a UML approach, in: B. Kleinjohann (Ed.), *Architecture and Design of Distributed Embedded Systems*, Kluwer Academic Publishers, Dordrecht, 2001, pp. 1–10.
- [58] R.J. Machado, J.M. Fernandes, A.J. Esteves, H.D. Santos, An evolutionary approach to the use of petri net based models: from parallel controllers to HW/SW co-design, in: A. Yakovlev, L. Gomes, L. Lavagno (Eds.), *Hardware Design and Petri Nets*, Kluwer Academic Publishers, Dordrecht, 2000, pp. 205–222.
- [59] J. Nützel, B. Däne, W. Fengler, Object nets for the design and verification of distributed and embedded applications, in: J.D.P. Rolim (Ed.), *Parallel and Distributed Processing, Proceedings of the 10th IPPS/SPDP'98 Workshops Held in Conjunction with the 12th International Parallel Processing Symposium and Ninth Symposium on Parallel and Distributed Processing*, Orlando, Florida, USA, March 30–April 3, 1998, Lecture Notes in Computer Science, Vol. 1388, Springer, Berlin, 1998, pp. 953–962.
- [60] J. Noguera, R. Badia, A HW/SW partitioning algorithm for dynamically reconfigurable architectures, in: *Proceedings of the Design, Automation and Test in Europe Conference (DATE'2001)*, Munich, Germany, March 12–16, 2001, pp. 729–734.
- [61] P. Garg, S.K. Shukla, R.K. Gupta, Efficient usage of concurrency models in an object-oriented co-design framework, in: *Proceedings of the Design Automation and Test in Europe (DATE'01)*, Designers Forum, Munich, Germany, March 2001.
- [62] C. Schulz-Key, T. Kuhn, W. Rosenstiel, A framework for system-level partitioning of object-oriented specifications, in: *Workshop on Synthesis and System Integration of Mixed Technologies (SASIMI'2001)*, Nara, Japan, 2001.

- [63] R.J. Machado, J.M. Fernandes, H.D. Santos, An object-oriented approach to the co-design of industrial control-based information systems, in: Fourth APCA Portuguese Conference on Automatic Control (CONTROLO 2000), Guimarães, Portugal, October, 2000.
- [64] H. Hallal, X. Kong, R. Negulescu, Experiments in modeling integrated circuit blocks by UML, Proceedings of the International Workshop on IP-based Circuits, 1999.
- [65] R. Chen, M. Sgroi, G. Martin, L. Lavagno, A.L. Sangiovanni-Vincentelli, J. Rabaey, UML and platform-based design, in: L. Lavagno, G. Martin, B. Selic (Eds.), UML for Real, Kluwer Academic Publishers, Dordrecht, 2003, pp. 107–126.
- [66] B. Liccardi, T. Maier-Komor, J.A. Oswald, M. Elkotob, G. Färber, A meta-modeling concept for embedded RT-systems design, in: 14th Euromicro Conference on Real-Time Systems, Vienna, Austria, 19–21 June 2002.
- [67] R. Damaševičius, V. Štuikys, Application of UML for hardware design based on design process model, Proceedings of the Asia South Pacific Design Automation Conference (ASP-DAC 2004), Yokohama, Japan, January 27–30, 2004, pp. 244–249.
- [68] W.E. McUmber, B.H.C. Cheng, UML-based analysis of embedded systems using a mapping to VHDL, Proceedings of the IEEE International Symposium on High Assurance Software Engineering (HASE'99), Washington, DC, USA, November 1999, pp. 56–63.
- [69] E. Gamma, R. Helm, R. Johnson, J. Vlissides, Design Patterns: Elements of Reusable Object-Oriented Software, Addison-Wesley, Reading, MA, 1995.
- [70] B.P. Douglass, Fine grained patterns for real-time systems, in: L. Lavagno, G. Martin, B. Selic (Eds.), UML for Real, Kluwer Academic Publishers, Boston, 2003, pp. 149–170.
- [71] B. Selic, Architectural Patterns for real-time systems, in: L. Lavagno, G. Martin, B. Selic (Eds.), UML for Real, Kluwer Academic Publishers, Boston, 2003, pp. 171–188.
- [72] N. Yoshida, Design patterns applied to object-oriented SoC design, in: 10th Workshop on Synthesis and System Integration of Mixed Technologies (SASIMI 2001), Nara, Japan, October 18–19, 2001.
- [73] P. Åström, S. Johansson, P. Nilsson, Application of software design patterns to DSP library design, Proceedings of the 14th International Symposium on System Synthesis (ISSS'01), Montreal, Canada, October 1–3, 2001, pp. 239–243.
- [74] R. Damaševičius, G. Majauskas, V. Štuikys, Application of design patterns for hardware design, Proceedings of the 40th Design Automation Conference (DAC 2003), Anaheim, CA, USA, June 2–6, 2003, pp. 48–53.
- [75] F. Doucet, R. K. Gupta, Microelectronic system-on-chip modeling using objects and their relationships, in: Online Symposium for Electrical Engineers (OSEE 2000), 2000.
- [76] G. Vanmeerbeeck, P. Schaumont, S. Vernalde, M. Engels, I. Bolsens, Hardware/software partitioning of embedded system in OC API-XL, in: Proceedings of the Ninth International Symposium on Hardware/Software Codesign (CODES'2001), Copenhagen, Denmark, April 25–27, 2001, pp. 30–35.
- [77] D. Kessner, Free-6502 core, 1999, <http://www.free-ip.com/6502/>.
- [78] LEOX Team, DRAGONFLY micro-core, 2001, <http://www.leox.org>.
- [79] Green Mountain Computing Systems, Inc., HC11 CPU core, 2000, <http://www.gmvhdl.com/hc11core.html>.
- [80] D. Wallner, AX8 core, 2001, <http://hem.passagen.se/dwallner/vhdl.html>.
- [81] T. Givargis, Intel 8051 micro-controller, 2000, <http://www.cs.ucr.edu/~dalton/i8051/i8051syn/>.
- [82] S. Kumar, A. Jantsch, J.-P. Soininen, M. Forsell, M. Millberg, J. Öberg, K. Tiensyrjä, A. Hemani, A network on chip architecture and design methodology, Proceedings of the IEEE Computer Society Annual Symposium on VLSI, April 2002, pp. 117–124.